

Programming with NUMA systems  
Where is data (Wally)?  
Is it out there?  
The truth is out there!



# Statement

Assume that the 128 elements of vectors a and b are distributed across the 4 nodes ( $M_{0-3}$ ) of a NUMA multiprocessor system as follows:

$M_0$	$M_1$	$M_2$	$M_3$
0..31	32..63	64..95	96..127

Each NUMA node consists of a single processor with its own local cache hierarchy and a portion of the physically distributed but logically shared main memory. Caches are kept coherent across (NUMA) nodes by using a directory-based coherence protocol. Each cache and memory line has a number of bits sufficient to store 4 consecutive elements of these vectors.

Given the following OpenMP parallel region:

```
#pragma omp parallel num_threads(4)
```

```
#pragma omp single
```

```
for (iter=0; i<99; iter++) {
```

```
    #pragma omp taskloop num_tasks(4)
```

```
    for (int i=0; i<128; i++)
```

```
        b[i] = foo1(a[i]);
```


```
    #pragma omp taskloop num_tasks(4)
```

```
    for (int i=0; i<128; i++)
```

```
        a[i] = foo2(b[i]);
```

```
}
```

- Explicit tasks
- Dynamic assignment to threads
- 32 iterations/task



- The execution of this parallel region does not cause coherence traffic; in other words, since each node has a portion of the vectors and the number of tasks generated in each taskloop equals the number of nodes, there will be no coherence commands interchanged between them.


## Let's analyze:

- 1) There is not control which data will be accessed by tasks in a taskloop
- 2) There is not control which data will be accessed by task between taskloops neither iterations

# Program Analysis: First Taskloop

- First taskloop creates 4 explicit tasks

```
#pragma omp parallel num_threads(4)
#pragma omp single
for (iter=0; i<99; iter++) {
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```



# Program Analysis: First Taskloop

```
#pragma omp parallel num_threads(4)
```

```
#pragma omp single
```

```
for (iter=0; i<99; iter++) {
```

```
    #pragma omp taskloop num_tasks(4)
```

```
    for (int i=0; i<128; i++)
```

```
        b[i] = foo1(a[i]);
```

```
    #pragma omp taskloop num_tasks(4)
```

```
    for (int i=0; i<128; i++)
```

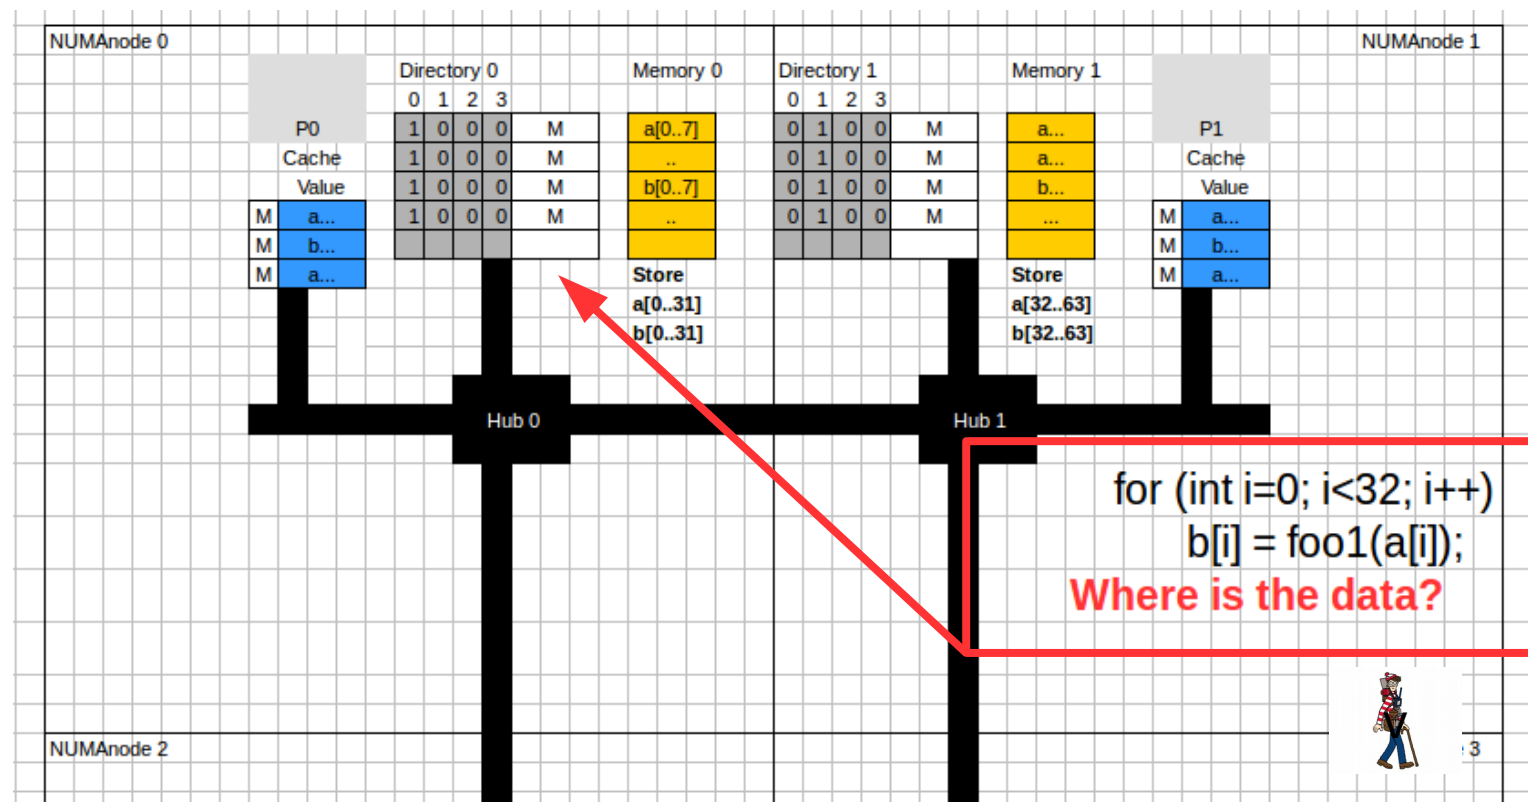
```
        a[i] = foo2(b[i]);
```

```
}
```

- First taskloop creates 4 explicit tasks

- Each task executes 32 iterations

- However we don't know which thread will run it
- In addition we don't know where this thread will be executed

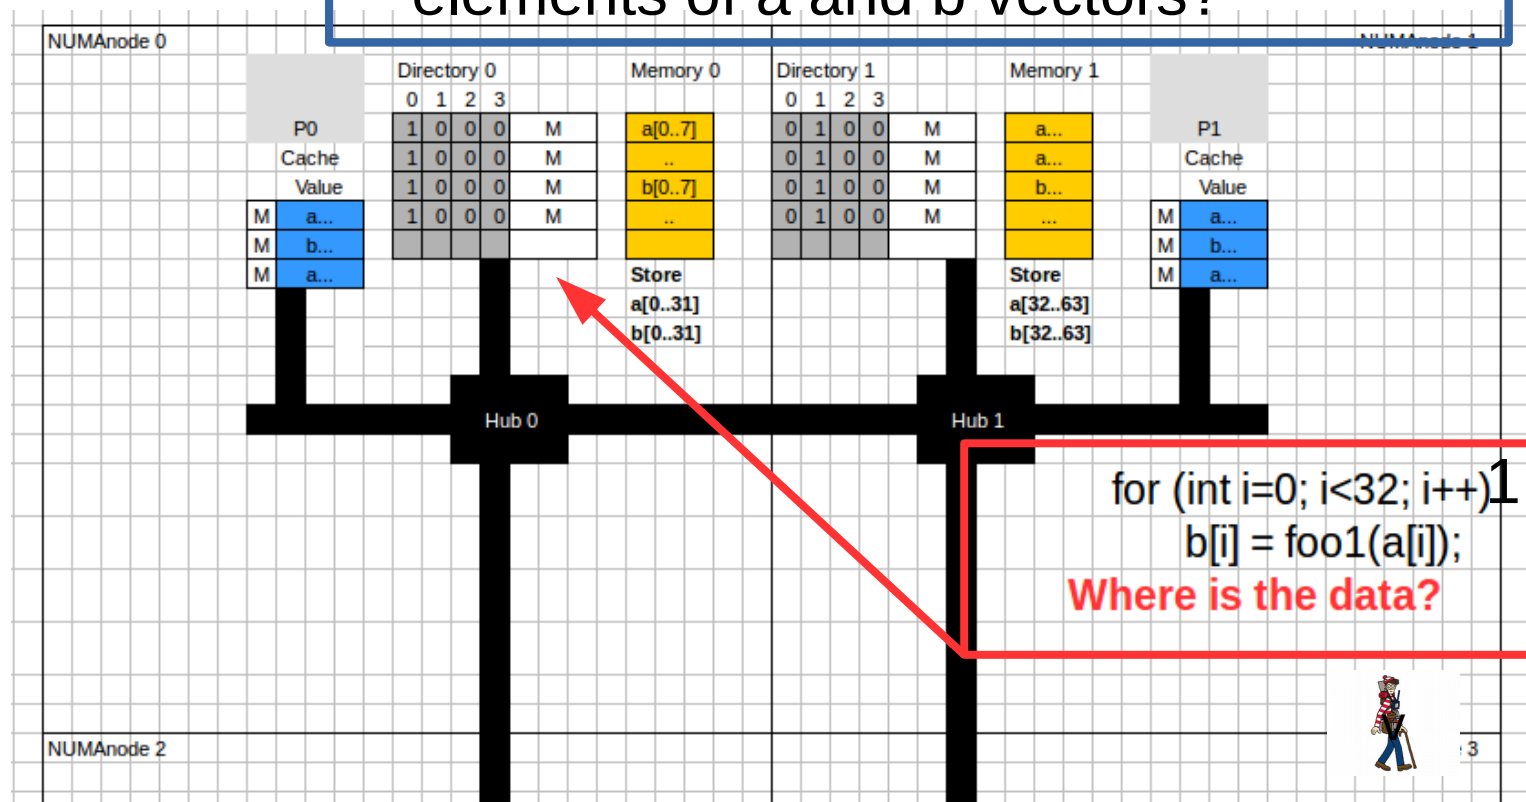


# Program Analysis: First Taskloop

```
#pragma omp parallel num_threads(4)
#pragma omp single
for (iter=0; i<99; iter++) {
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

- First taskloop creates 4 explicit tasks
- Each task has 32 iterations assigned
- Taskloop has an implicit taskwait

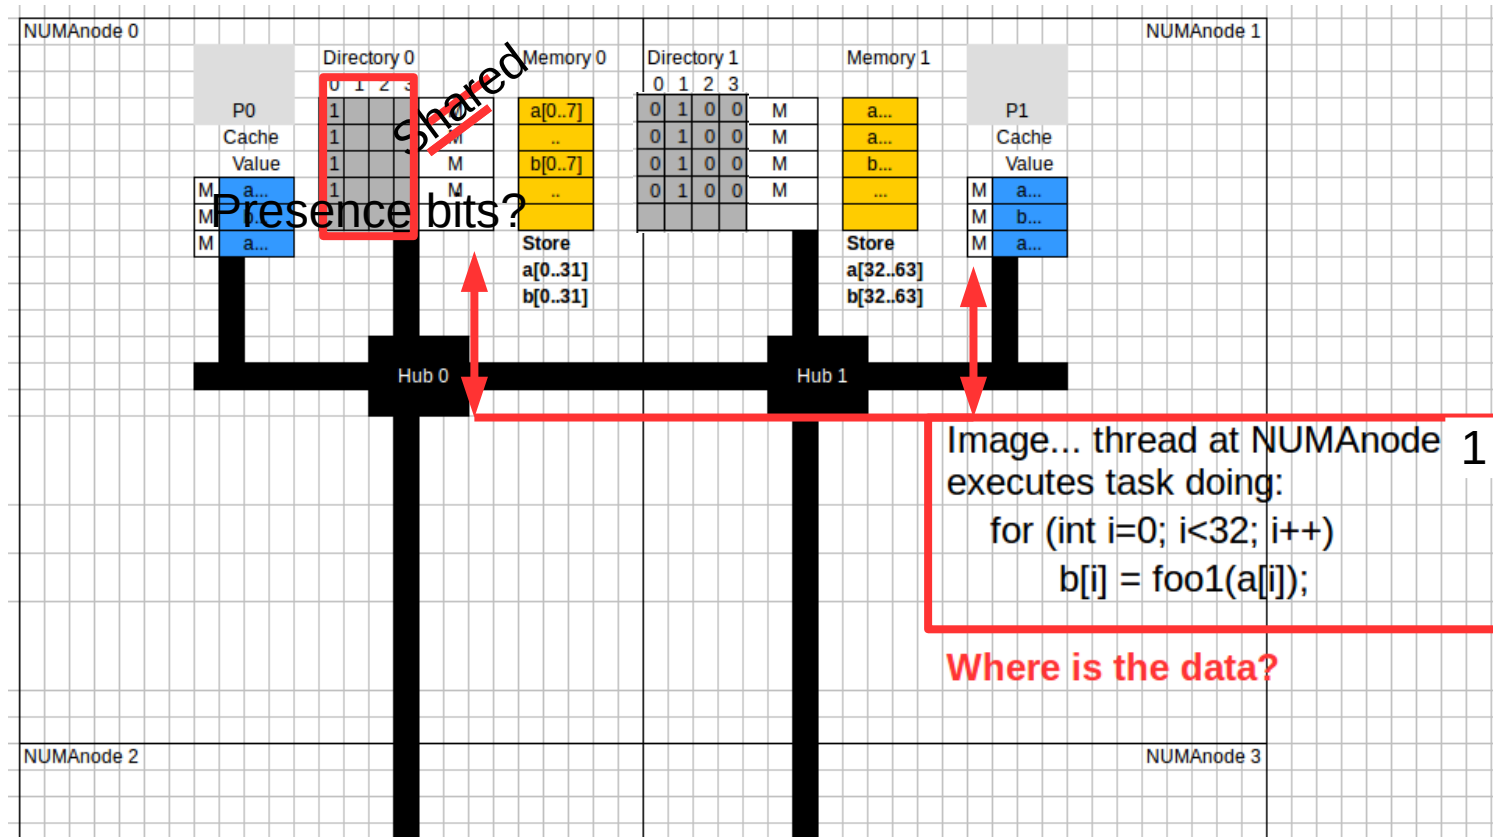
- Who perform the access to the first 32 elements of a and b vectors?



```
for (int i=0; i<32; i++)1
    b[i] = foo1(a[i]);
Where is the data?
```

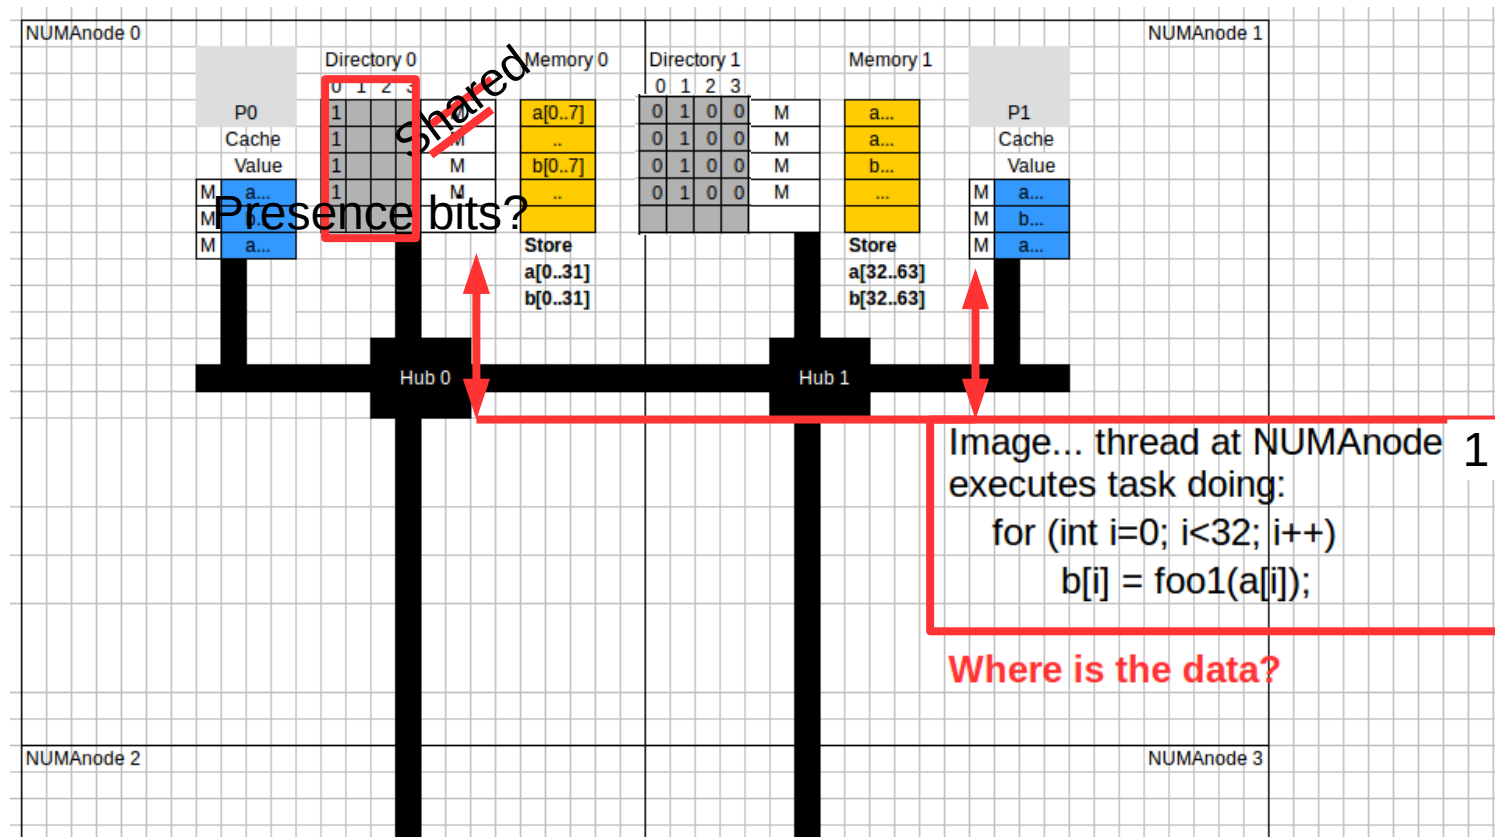


# Program Analysis: First Taskloop





# Program Analysis: First Taskloop



Let's see the detail

# Coherence data after first taskloop




Image... thread at NUMA Node 1 executes task doing:

```
for (int i=0; i<32; i++)
```

```
    b[i] = foo1(a[i]);
```

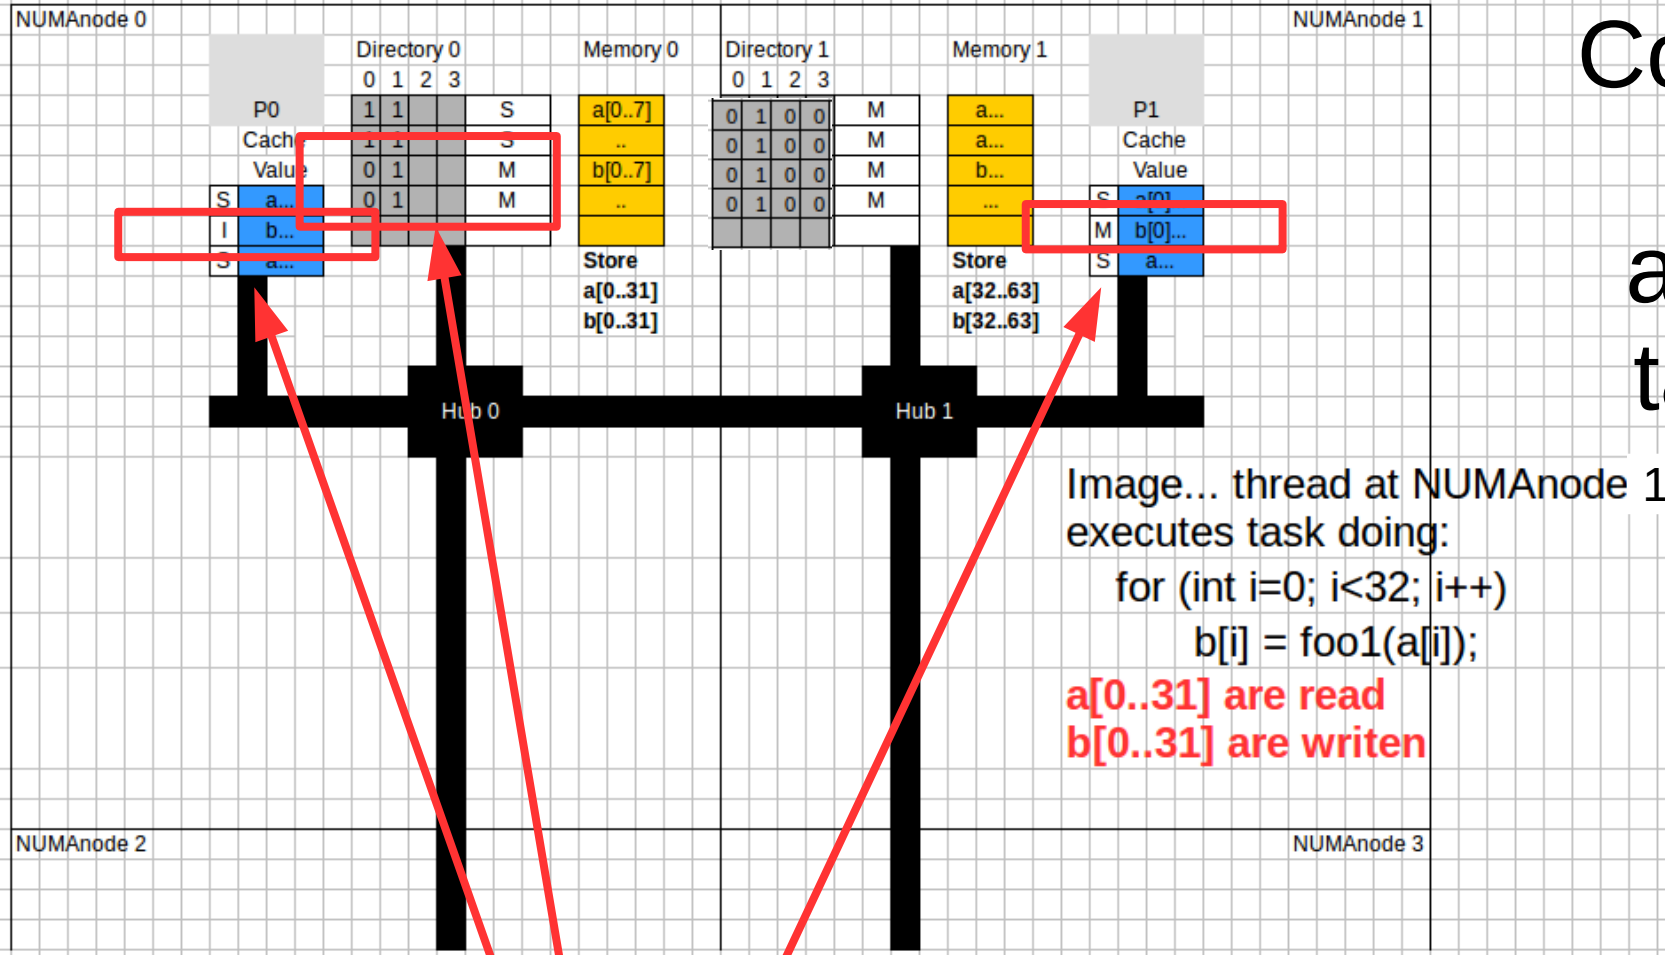
**a[0..31] are read**  
**b[0..31] are written**

NUMA Node home is NUMA Node 0

First taskloop:

- a[0..31] was read at NUMA Node 1
  - Directory indicates copies at caches of NumA Node 0 and 1
  - Cache lines at P1 has shared copies

# Coherence data after first taskloop

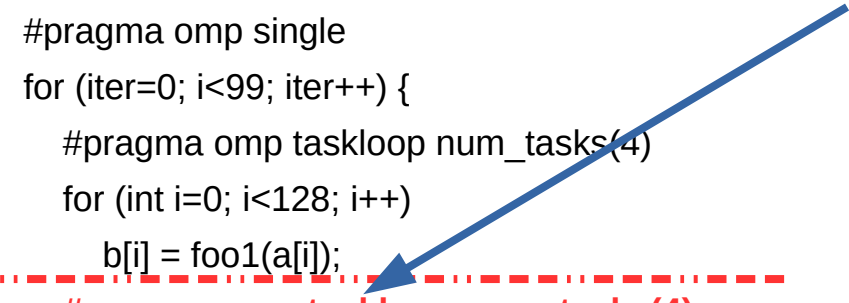


NUMA node home is NUMA node 0

- a[0..31] was read at NUMA node 1
  - Directory indicates copies at caches of NumaNode0 and 1
  - Cache lines at P1 has shared copies
- b[0..31] was modified at NUMA node 1
  - Directory indicates modified copies at NumaNode 1
  - Cache lines at P1 has modified copies

# Program Analysis: Second Taskloop

```
#pragma omp parallel num_threads(4)
#pragma omp single
for (iter=0; i<99; iter++) {
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```



- Second Taskloop creates 4 tasks
- Each task has 32 iterations assigned
- Are the same iterations (portion of vectors a and b) assigned to the same thread?

# Program Analysis: Second Taskloop


```
#pragma omp parallel num_threads(4)
#pragma omp single
for (iter=0; i<99; iter++) {
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

- Second Taskloop creates 4 tasks
- Each task has 32 iterations assigned

- Are the same iterations assigned to the same thread? **NOOOO.....**

Initially, after first taskloop:

- NUMA node home is NUMA node 0
- a[0..31] was read at NUMA node 1
  - Memory lines have copies at caches of P0 and P1
- b[0..31] was modified at NUMA node 1
  - Memory lines only have modified copies at P1



Assume now that a thread running at NUMA node 2 or 3 does task accessing a[0..31] and b[0..31] ...

- Invalidations have to be done
- Presence bits has to be updated



- The execution of this parallel region does not cause coherence traffic; in other words, since each node has a portion of the vectors and the number of tasks generated in each taskloop equals the number of nodes, there will be no coherence commands interchanged between them.



The execution of this parallel region does not cause coherence traffic; in other words, since each node has a portion of the vectors and the number of tasks generated in each taskloop equals the number of nodes, there will be no coherence commands interchanged between them.

- The execution of this parallel region only causes coherence traffic during the execution of the first iteration (iter=0) of the outer loop; in other words, once the elements are brought into cache there will be no coherence commands interchanged between nodes.

```
#pragma omp parallel num_threads(4)
#pragma omp single
for (iter=0; i<99; iter++) {
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

- Taskloops will create new tasks
- Tasks are dynamically assigned





The execution of this parallel region only causes coherence traffic during the execution of the first iteration (iter=0) of the outer loop; in other words, once the elements are brought into cache there will be no coherence commands interchanged between nodes.

```
#pragma omp parallel num_threads(4)
#pragma omp single
for (iter=0; i<99; iter++) {
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

- Taskloops will create new tasks
- Tasks are dynamically assigned
- From one iteration to another iteration there may be different assignment
  - Same problem as before

- The execution of this parallel region may cause coherence traffic during the execution of all iterations of the outer loop; in other words, it is highly probable that the processor in a node accesses (reads and/or writes) data cached in the cache of other nodes.

```
#pragma omp parallel num_threads(4)
#pragma omp single
for (iter=0; i<99; iter++) {
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

- The execution of this parallel region may cause coherence traffic during the execution of all iterations of the outer loop; in other words, it is highly probable that the processor in a node accesses (reads and/or writes) data cached in the cache of other nodes.

```
#pragma omp parallel num_threads(4)
#pragma omp single
for (iter=0; i<99; iter++) {
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

- There may be plenty of coherence traffic due to accesses to data cached in the cache of other nodes



The execution of this parallel region may cause coherence traffic during the execution of all iterations of the outer loop; in other words, it is highly probable that the processor in a node accesses (reads and/or writes) data cached in the cache of other nodes.

```
#pragma omp parallel num_threads(4)
#pragma omp single
for (iter=0; i<99; iter++) {
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

- There may be plenty of coherence traffic due to accesses to data cached in the cache of other nodes

# Statement

Assume that the 128 elements of vectors *a* and *b* are distributed across the 4 nodes ( $M_{0-3}$ ) of a NUMA multiprocessor system as follows:

$M_0$	$M_1$	$M_2$	$M_3$
0..31	32..63	64..95	96..127

Each NUMA node consists of a single processor with its own local cache hierarchy and a portion of the physically distributed but logically shared main memory. Caches are kept coherent across (NUMA) nodes by using a directory-based coherence protocol. Each cache and memory line has a number of bits sufficient to store 4 consecutive elements of these vectors.

For the same architecture, if the parallel region is changed as follows:

```
#pragma omp parallel num_threads(4)
for (iter=0; i<99; iter++) {
    #pragma omp for schedule(static)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp for schedule(static)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

# Statement

Assume that the 128 elements of vectors a and b are distributed across the 4 nodes ( $M_{0-3}$ ) of a NUMA multiprocessor system as follows:

$M_0$	$M_1$	$M_2$	$M_3$
0..31	32..63	64..95	96..127

Each NUMA node consists of a single processor with its own local cache hierarchy and a portion of the physically distributed but logically shared main memory. Caches are kept coherent across (NUMA) nodes by using a directory-based coherence protocol. Each cache and memory line has a number of bits sufficient to store 4 consecutive elements of these vectors.

For the same architecture, if the parallel region is changed as follows:

```
#pragma omp parallel num_threads(4)
```

```
for (iter=0; i<99; iter++) {
```

```
    #pragma omp for schedule(static)
```

```
    for (int i=0; i<128; i++)
```

```
        b[i] = foo1(a[i]);
```

```
    #pragma omp for schedule(static)
```

```
    for (int i=0; i<128; i++)
```

```
        a[i] = foo2(b[i]);
```

```
}
```

- Implicit tasks
- Schedule static (N/threads)



# Program Analysis

```
#pragma omp parallel num_threads(4)
for (iter=0; i<99; iter++) {
    #pragma omp for schedule(static)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp for schedule(static)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

• Schedule static assignning (N/4) to each thread

- Thread 0: b[0..31],a[0..31]
- Thread 1: b[32..63],a[32..63]
- Thread 2: b[64..95],a[64..95]
- Thread 3: b[96..127],a[96..127]



# Program Analysis

```
#pragma omp parallel num_threads(4)
for (iter=0; i<99; iter++) {
    #pragma omp for schedule(static)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp for schedule(static)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

- Schedule static assignning (N/4) to each thread
- There is an implicit barrier at the end of each omp for
- Each omp for has the same schedule

# Program Analysis

```
#pragma omp parallel num_threads(4)
for (iter=0; i<99; iter++) {
    #pragma omp for schedule(static)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp for schedule(static)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

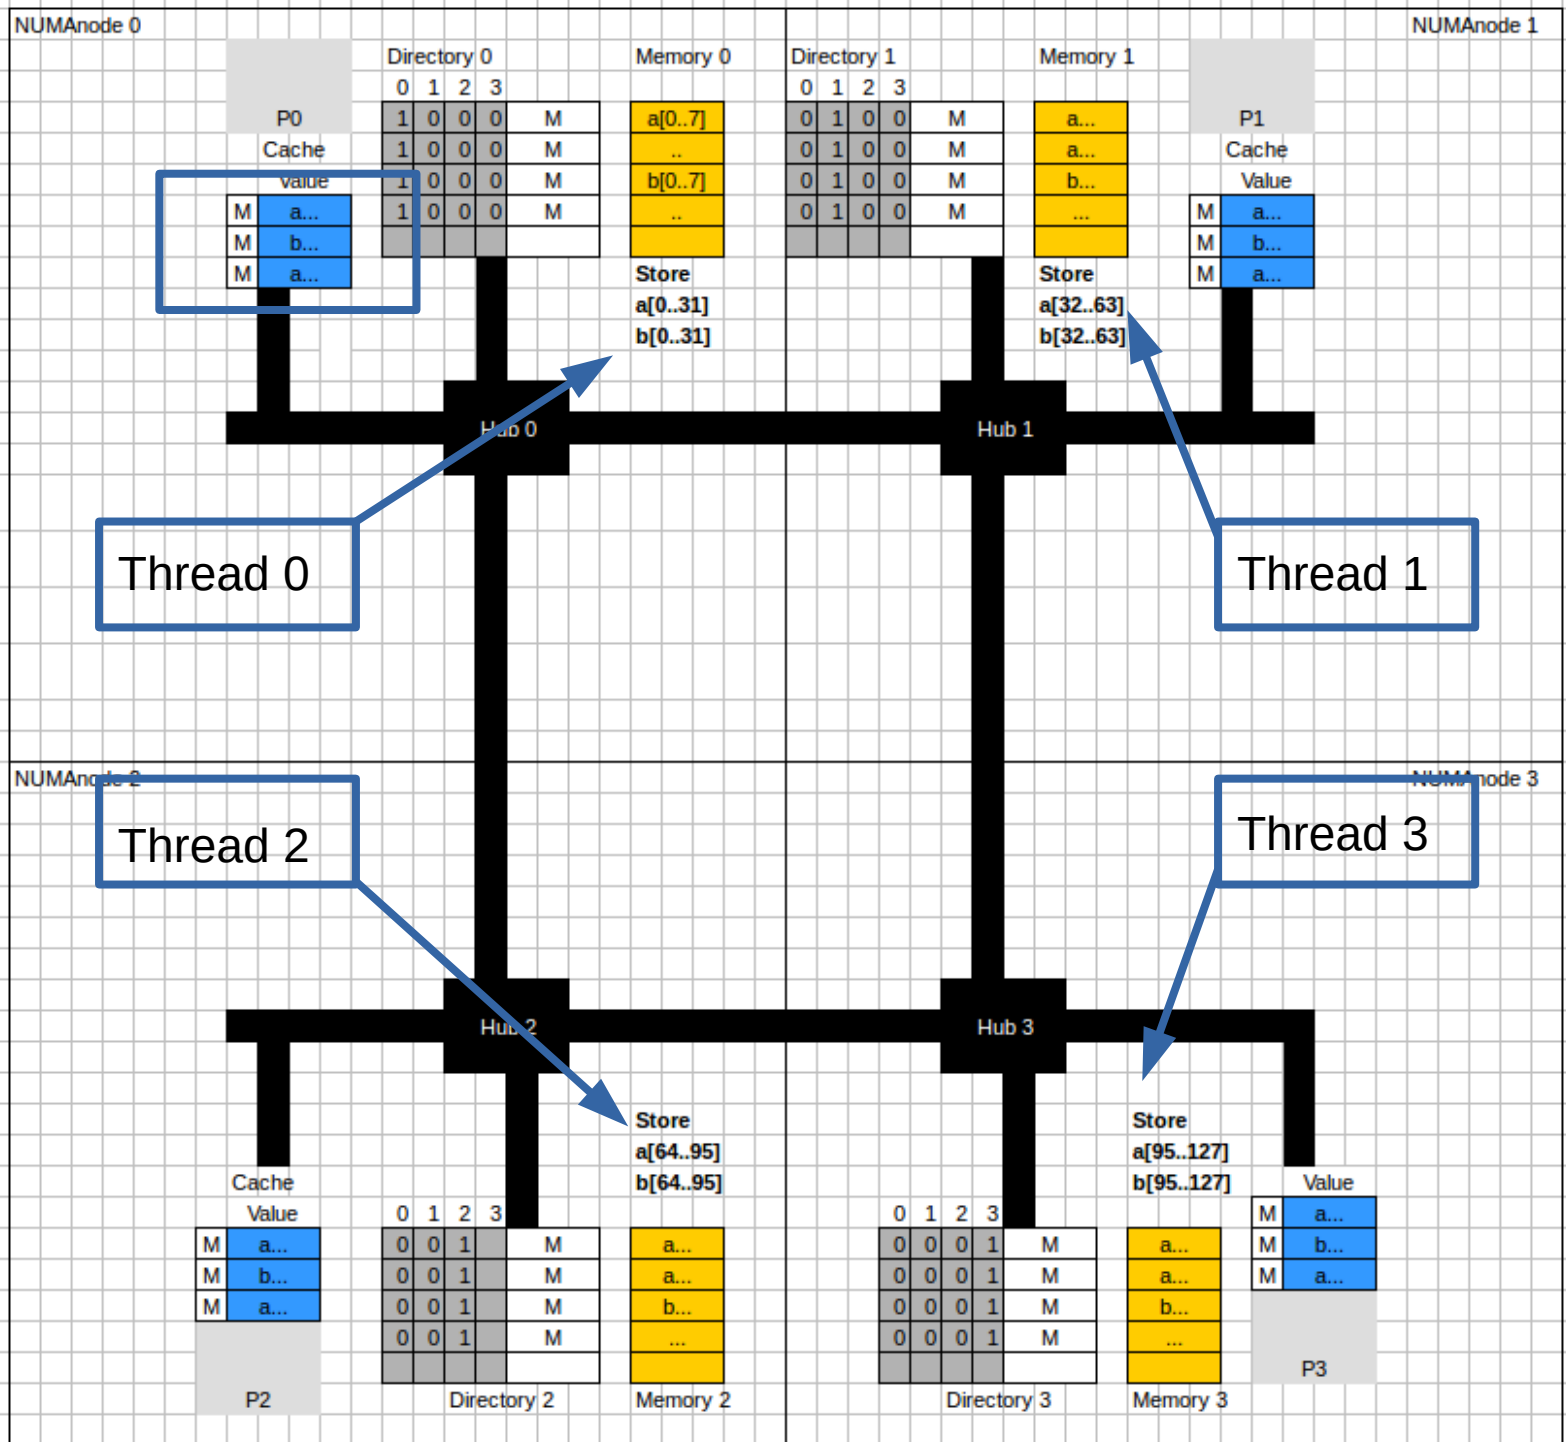
- Schedule static assignming (N/4) to each thread

- There is an implicit barrier at the end of each omp for

- **Assume thread i runs in NumaNode/Processor i and..**

- Thread 0: b[0..31],a[0..31]
- Thread 1: b[32..63],a[32..63]
- Thread 2: b[64..95],a[64..95]
- Thread 3: b[96..127],a[96..127]

**Where is the data that is accessed by each thread?**



- The execution of this parallel region does not cause coherence traffic because each node has a portion of the vectors and the static assignment will always assign the same iterations to each thread. Thus, each thread will only access the portion of the vectors stored within the memory of a node. Consequently, there will be no coherence commands interchanged between nodes
- The execution of this parallel region only causes coherence traffic during the execution of the first iteration ( $\text{iter}=0$ ) of the outer loop; in other words, once the elements are brought into cache there will be no coherence commands interchanged between nodes.
- The execution of this parallel region may cause coherence traffic during the execution of all iterations of the outer loop; in other words, it is highly probable that the processor in a node accesses (reads and/or writes) data cached in the cache of other nodes.

- The execution of this parallel region does not cause coherence traffic because each node has a portion of the vectors and the static assignment will always assign the same iterations to each thread. Thus, each thread will only access the portion of the vectors stored within the memory of a node. Consequently, there will be no coherence commands interchanged between nodes
- The execution of this parallel region only causes coherence traffic during the execution of the first iteration ( $\text{iter}=0$ ) of the outer loop; in other words, once the elements are brought into cache there will be no coherence commands interchanged between nodes.



The execution of this parallel region may cause coherence traffic during the execution of all iterations of the outer loop; in other words, it is highly probable that the processor in a node accesses (reads and/or writes) data cached in the cache of other nodes.

- The execution of this parallel region does not cause coherence traffic because each node has a portion of the vectors and the static assignment will always assign the same iterations to each thread. Thus, each thread will only access the portion of the vectors stored within the memory of a node. Consequently, there will be no coherence commands interchanged between nodes



The execution of this parallel region only causes coherence traffic during the execution of the first iteration ( $\text{iter}=0$ ) of the outer loop; in other words, once the elements are brought into cache there will be no coherence commands interchanged between nodes.



The execution of this parallel region may cause coherence traffic during the execution of all iterations of the outer loop; in other words, it is highly probable that the processor in a node accesses (reads and/or writes) data cached in the cache of other nodes.



The execution of this parallel region does not cause coherence traffic because each node has a portion of the vectors and the static assignment will always assign the same iterations to each thread. Thus, each thread will only access the portion of the vectors stored within the memory of a node. Consequently, there will be no coherence commands interchanged between nodes



The execution of this parallel region only causes coherence traffic during the execution of the first iteration ( $\text{iter}=0$ ) of the outer loop; in other words, once the elements are brought into cache there will be no coherence commands interchanged between nodes.



The execution of this parallel region may cause coherence traffic during the execution of all iterations of the outer loop; in other words, it is highly probable that the processor in a node accesses (reads and/or writes) data cached in the cache of other nodes.





The execution of this parallel region does not cause coherence traffic because each node has a portion of the vectors and the static assignment will always assign the same iterations to each thread. Thus, each thread will only access the portion of the vectors stored within the memory of a node. Consequently, there will be no coherence commands interchanged between nodes



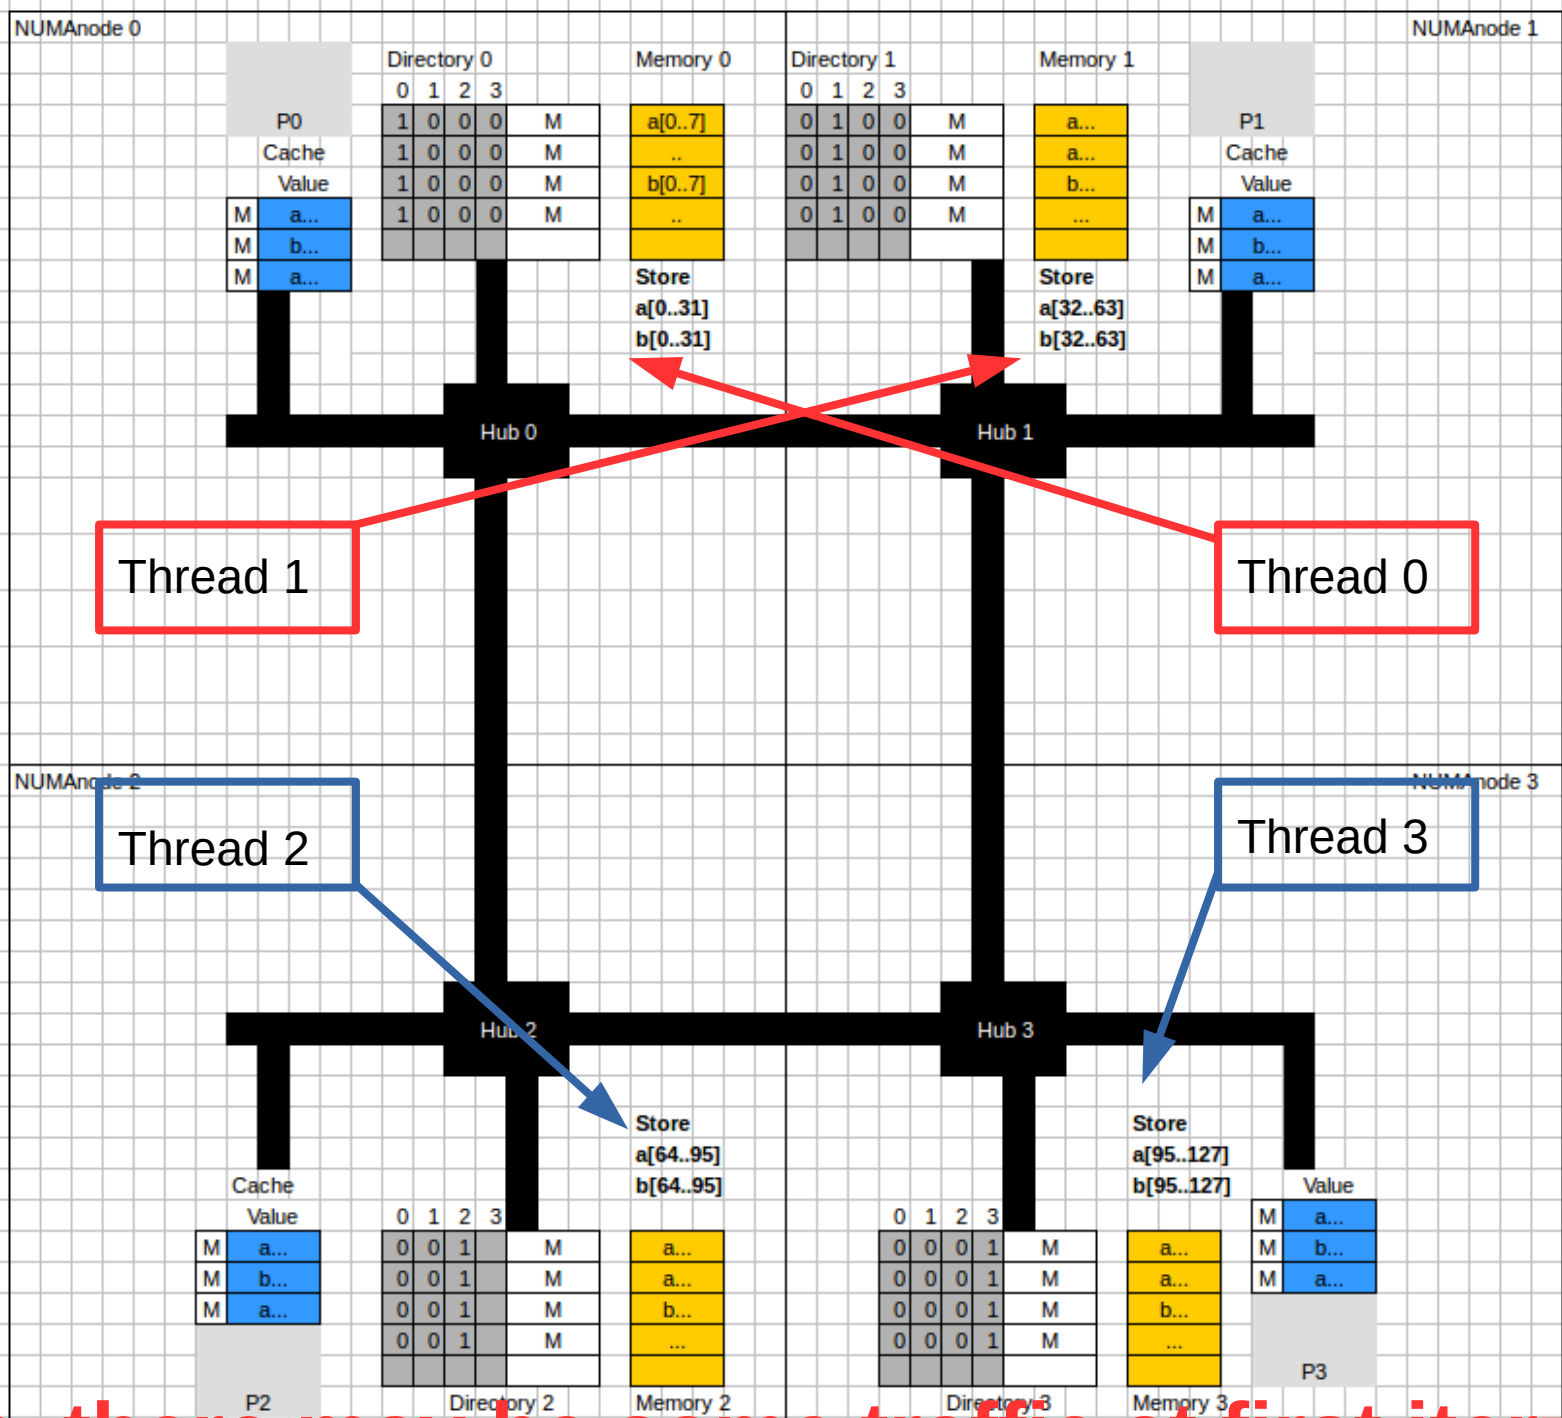
The execution of this parallel region only causes coherence traffic during the execution of the first iteration (iter=0) of the outer loop; in other words, once the elements are brought into cache there will be no coherence commands interchanged between nodes.



The execution of this parallel region may cause coherence traffic during the execution of all iterations of the outer loop; in other words, it is highly probable that the processor in a node accesses (reads and/or writes) data cached in the cache of other nodes.

**If we don't assume that thread i goes to NumaNode/Processor i... would be the answers correct?**





Then, there may be some traffic at first iter=0

# Program Analysis

```
#pragma omp parallel num_threads(4)
for (iter=0; i<99; iter++) {
    #pragma omp for schedule(static)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp for schedule(static)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

- **First iteration iter=0**

- thread 0 and thread 1 would need to access NumaNode 1 and 0 respectively and invalidate copies of the caches of the another processor, and update presence bits.
- Thread 2 and thread 3 would locally access their local data.

- **Iter 1 and following**

- Caches of each processor will have already local copies of the data to be accessed in iter 1 and following
- No more traffic is necessary

- The execution of this parallel region does not cause coherence traffic because each node has a portion of the vectors and the static assignment will always assign the same iterations to each thread. Thus, each thread will only access the portion of the vectors stored within the memory of a node. Consequently, there will be no coherence commands interchanged between nodes
- The execution of this parallel region only causes coherence traffic during the execution of the first iteration (iter=0) of the outer loop; in other words, once the elements are brought into cache there will be no coherence commands interchanged between nodes.
- The execution of this parallel region may cause coherence traffic during the execution of all iterations of the outer loop; in other words, it is highly probable that the processor in a node accesses (reads and/or writes) data cached in the cache of other nodes.

# Program Analysis

```
#pragma omp parallel num_threads(4)
for (iter=0; i<99; iter++) {
    #pragma omp for schedule(static)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp for schedule(static)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

- **First iteration iter=0**

- thread 0 and thread 1 would need to access NumaNode 1 and 0 respectively and invalidate copies of the caches of the another processor, and update presence bits.
- Thread 2 and thread 3 would locally access their local data.

- **Iter 1 and following**

- Caches of each processor will have already local copies of the data to be accessed in iter 1 and following
- No more traffic is necessary



The execution of this parallel region does not cause coherence traffic because each node has a portion of the vectors and the static assignment will always assign the same iterations to each thread. Thus, each thread will only access the portion of the vectors stored within the memory of a node. Consequently, there will be no coherence commands interchanged between nodes



The execution of this parallel region only causes coherence traffic during the execution of the first iteration (iter=0) of the outer loop; in other words, once the elements are brought into cache there will be no coherence commands interchanged between nodes.



The execution of this parallel region may cause coherence traffic during the execution of all iterations of the outer loop; in other words, it is highly probable that the processor in a node accesses (reads and/or writes) data cached in the cache of other nodes.

# Statement

Assume now that the 128 elements of vectors a and b are allocated only in the node associated to  $M_0$ . This has happened because 1) the programmer simply forgot to parallelize the initialization loop:

```
for (int i=0; i<128; i++) {  
    a[i] = random();  
    b[i] = random();  
}
```

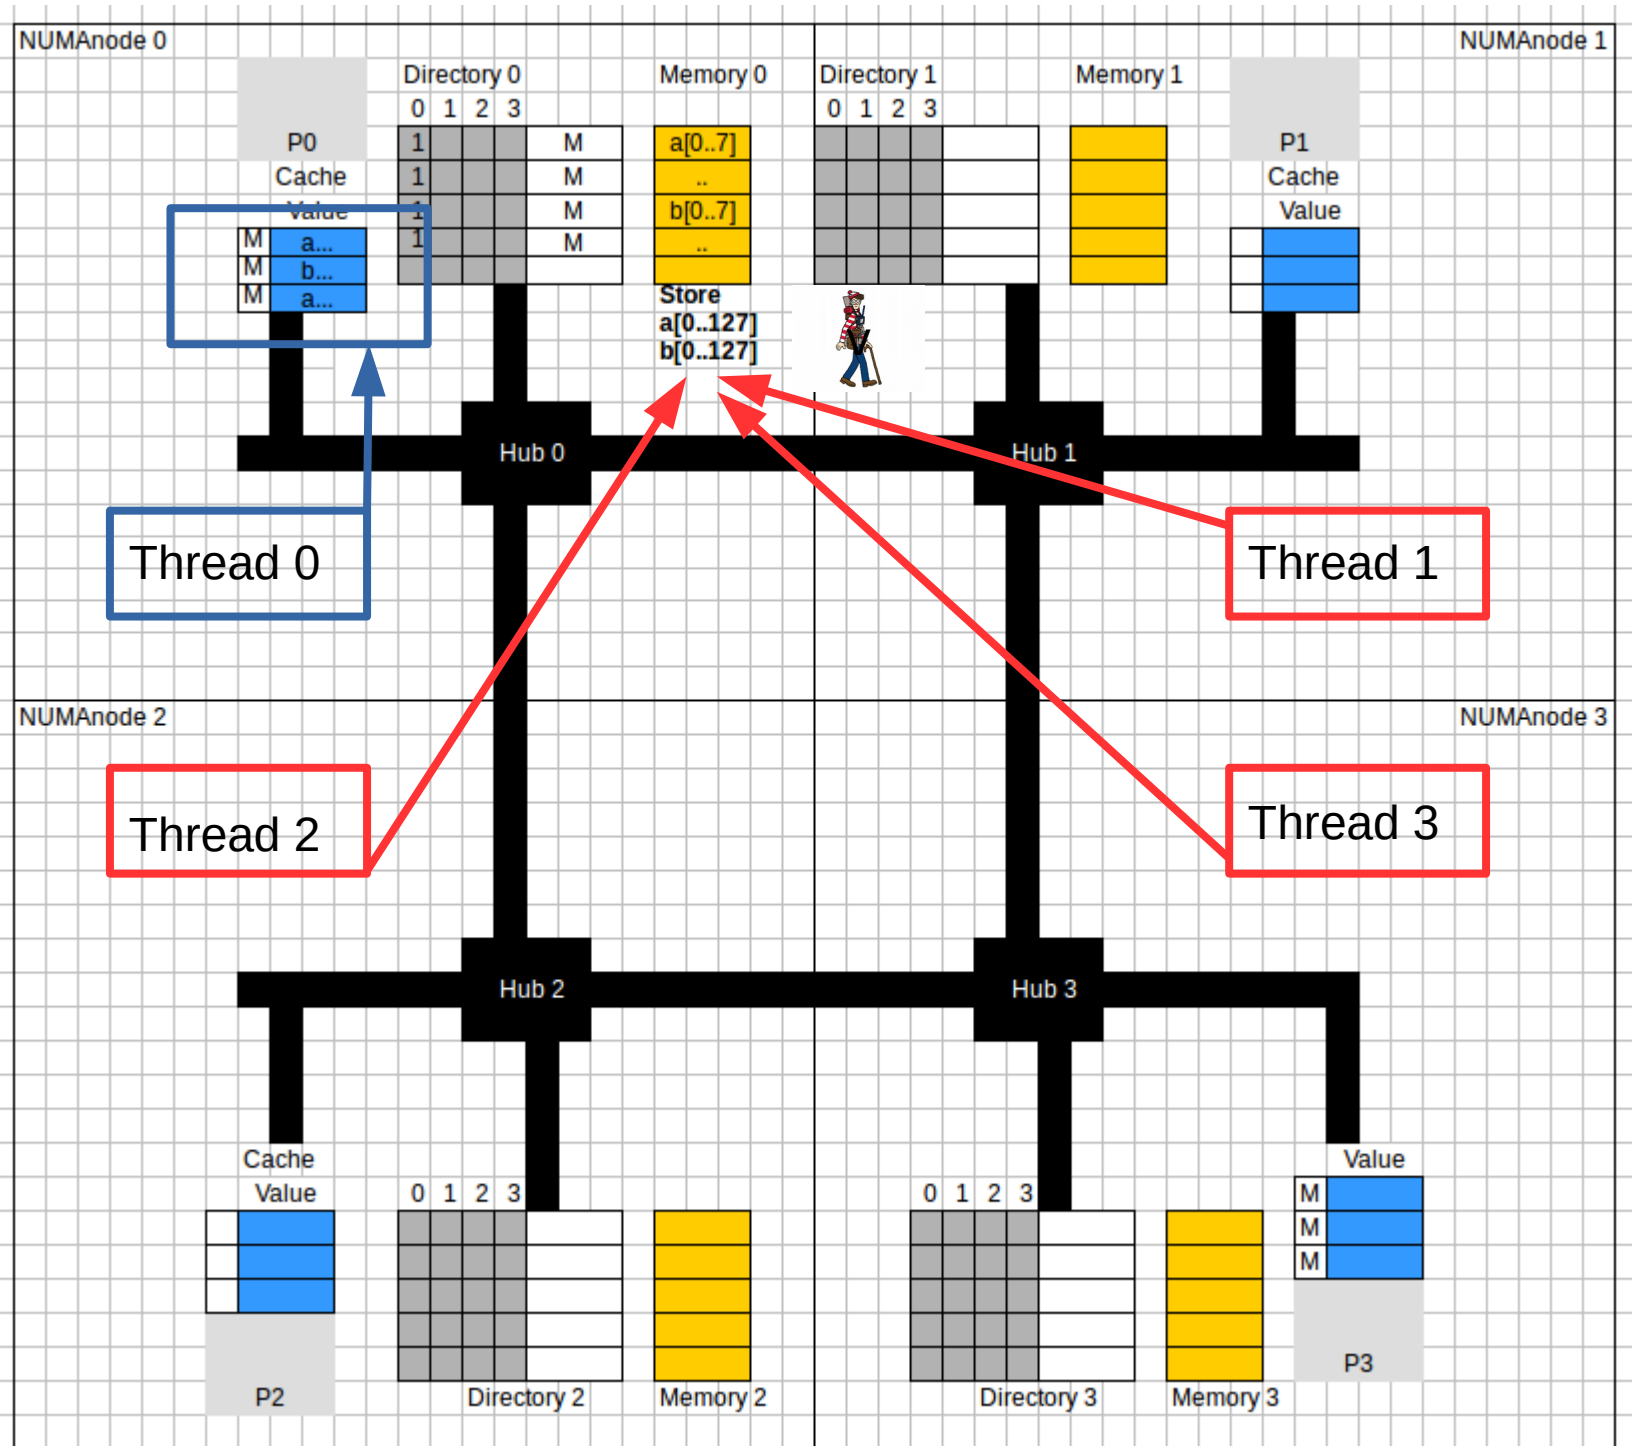
and 2) because of the first touch policy explained in class. Compared to the previous two questions, in which vector elements were distributed among all the nodes, which of the following statements is true:

# Program Analysis

```
#pragma omp parallel num_threads(4)
for (iter=0; i<99; iter++) {
    #pragma omp for schedule(static)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp for schedule(static)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

```
#pragma omp parallel num_threads(4)
#pragma omp single
for (iter=0; i<99; iter++) {
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

- ❑ The two previous versions of the parallel region (the one with taskloop and the other with static schedule) will not cause now any coherence traffic since vectors are stored in a single node ( $M_0$ ) of the system.
- ❑ The version based on taskloop will now behave even worse since all coherence commands should be served by the same directory structure in node  $M_0$ .
- ❑ The version based on static schedule will have a similar behaviour after the execution of the first iteration ( $\text{iter}=0$ ) of the outer loop. However, during the first iteration of the outer loop the behaviour will be worse now since all coherence commands should be served by the same directory structure in node  $M_0$ .



# Program Analysis

```
#pragma omp parallel num_threads(4)
for (iter=0; i<99; iter++) {
    #pragma omp for schedule(static)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp for schedule(static)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

```
#pragma omp parallel num_threads(4)
#pragma omp single
for (iter=0; i<99; iter++) {
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```



The two previous versions of the parallel region (the one with taskloop and the other with static schedule) will not cause now any coherence traffic since vectors are stored in a single node ( $M_0$ ) of the system.

- ☐ The version based on taskloop will now behave even worse since all coherence commands should be served by the same directory structure in node  $M_0$ .
- ☐ The version based on static schedule will have a similar behaviour after the execution of the first iteration ( $\text{iter}=0$ ) of the outer loop. However, during the first iteration of the outer loop the behaviour will be worse now since all coherence commands should be served by the same directory structure in node  $M_0$ .

# Program Analysis

```
#pragma omp parallel num_threads(4)
for (iter=0; i<99; iter++) {
    #pragma omp for schedule(static)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp for schedule(static)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

- **First iteration iter=0**

- Thread 1,2,3 will need to access NumaNode 0 to ask for data, update presence bits, and notify any WrReq or UpgrReq
- Thread 0 will locally access their local data.

- **Iter 1 and following**

- Caches of each processor will have already local copies of the data to be accessed in iter 1 and following
- No more traffic is necessary



The two previous versions of the parallel region (the one with taskloop and the other with static schedule) will not cause now any coherence traffic since vectors are stored in a single node ( $M_0$ ) of the system.

- ☐ The version based on taskloop will now behave even worse since all coherence commands should be served by the same directory structure in node  $M_0$ .



The version based on static schedule will have a similar behaviour after the execution of the first iteration (iter=0) of the outer loop. However, during the first iteration of the outer loop the behaviour will be worse now since all coherence commands should be served by the same directory structure in node  $M_0$ .



# Program Analysis

```
#pragma omp parallel num_threads(4)
#pragma omp single
for (iter=0; i<99; iter++) {
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

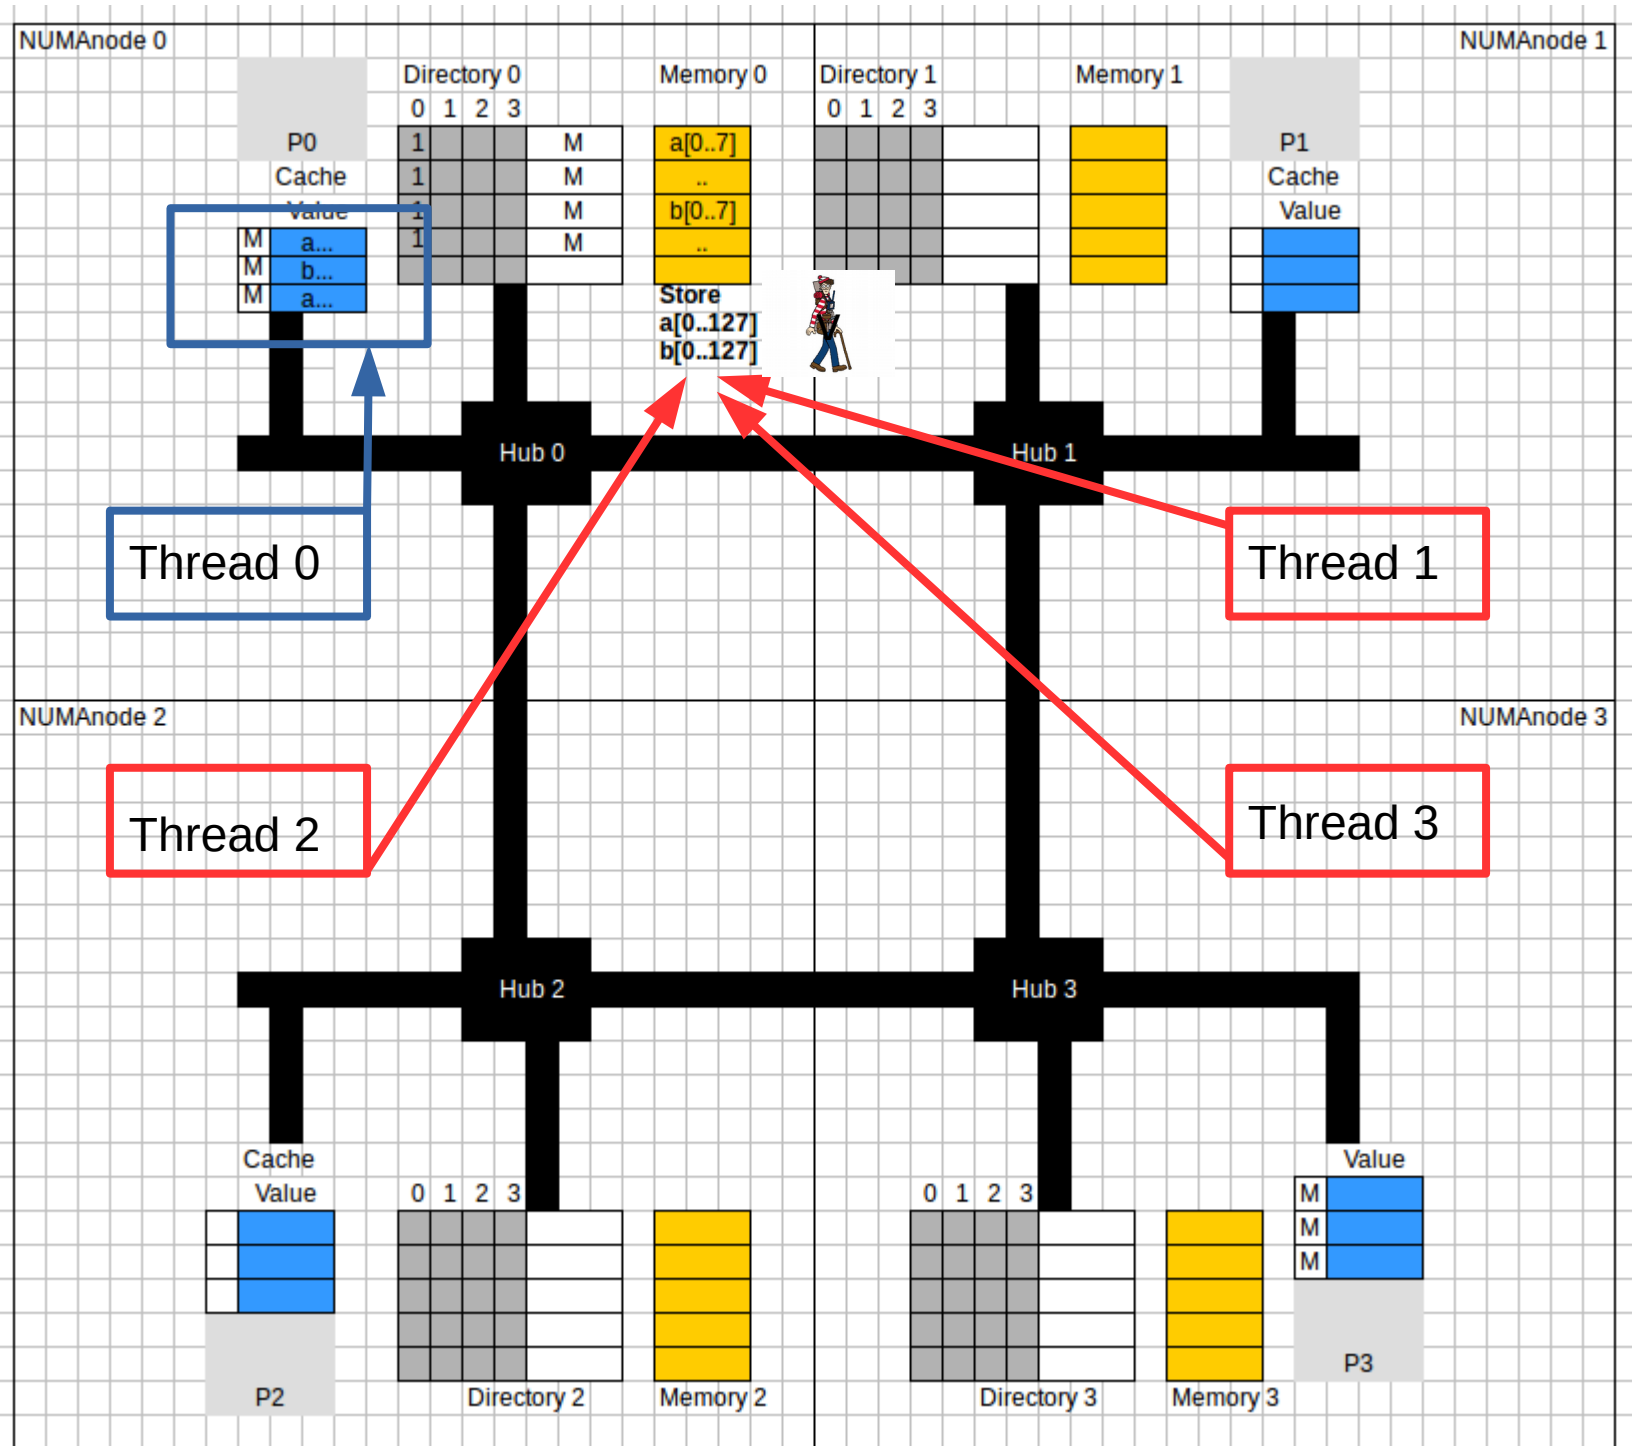


The two previous versions of the parallel region (the one with taskloop and the other with static schedule) will not cause now any coherence traffic since vectors are stored in a single node ( $M_0$ ) of the system.

- ☐ The version based on taskloop will now behave even worse since all coherence commands should be served by the same directory structure in node  $M_0$ .



The version based on static schedule will have a similar behaviour after the execution of the first iteration (iter=0) of the outer loop. However, during the first iteration of the outer loop the behaviour will be worse now since all coherence commands should be served by the same directory structure in node  $M_0$ .



# Program Analysis

```
#pragma omp parallel num_threads(4)
#pragma omp single
for (iter=0; i<99; iter++) {
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        b[i] = foo1(a[i]);
    #pragma omp taskloop num_tasks(4)
    for (int i=0; i<128; i++)
        a[i] = foo2(b[i]);
}
```

- Now all threads running tasks will access to the same NumaNode: NumaNode 0



The two previous versions of the parallel region (the one with taskloop and the other with static schedule) will not cause now any coherence traffic since vectors are stored in a single node ( $M_0$ ) of the system.



The version based on taskloop will now behave even worse since all coherence commands should be served by the same directory structure in node  $M_0$ .



The version based on static schedule will have a similar behaviour after the execution of the first iteration (iter=0) of the outer loop. However, during the first iteration of the outer loop the behaviour will be worse now since all coherence commands should be served by the same directory structure in node  $M_0$ .