

Dado el siguiente código secuencial en C que encuentra en un vector DB la primera posición en la que aparece una determinada clave key:

```
int main() {  
    double key = 1.25;  
    double * DB = (double *) malloc(sizeof(double) * DBsize);  
    initialize(DB, &DBsize); // initialize elements in DB  
    unsigned long position = DBsize;  
    for (unsigned long i = 0; (i < DBsize) && (position == DBsize); i++)  
        if (DB[i] == key) position = i;  
}
```

Y la siguiente solución incompleta para la paralelización del bucle for:

```
#pragma omp parallel  
{  
    unsigned long i, num_elems, lower;  
    for (i = lower; (i < (lower + num_elems)) && (i < position); i++) {  
        #pragma omp critical  
        if ((DB[i] == key) && (i < position)) position = i;  
    }  
}
```


```
#pragma omp parallel
{
    unsigned long i, BS, lower;
    for (i = lower; (i < (lower + BS)) && (i < position); i++) {
        #pragma omp critical
            if ((DB[i] == key) && (i < position)) position = i;
    }
}
```

Se pide:

- Completad la solución incompleta anterior (con las sentencias y declaraciones de datos necesarias) para que cumpla con las siguientes condiciones, sin preocuparse por posibles problemas de rendimiento: 1) el reparto de iteraciones a procesadores obedezca a una descomposición de datos geométrica tipo BLOCK (es decir, a cada procesador se le asocian $DBsize/P$ elementos consecutivos, siendo P el número de procesadores); 2) P no tiene por qué dividir de forma entera $DBsize$, en cuyo caso se deberá maximizar el balanceo de carga; 3) la solución debe permitir que un procesador finalice su ejecución tan pronto encuentre `key` o detecte que no contribuirá a la solución final; y 4) no puede utilizarse el `#pragma omp for` para realizar el reparto de iteraciones.
- Modificad el código anterior para que se mejore el rendimiento de forma substancial, reduciendo al mínimo la secuencialización que introduce la sincronización actual.
- Proponed e implementad una descomposición de datos geométrica alternativa que reduzca el tiempo de ejecución requerido, en media, para encontrar la primera posición en la que aparece la clave `key`.

a) Input BLOCK Geometric Data Decomposition:

- Block size = DBsize / P
- $(\text{DBsize} \% P)$ can be > 0
- Execution at each processor must end as soon as possible
- You cannot use *omp for*




```
#pragma omp parallel
{
    int myid = omp_get_thread_num ();
    int P = omp_get_num_threads ();
    ...
    unsigned long lower =
    unsigned long upper =
    ...

    unsigned long i;
    for (i = lower; ( ... ); i++) {
        #pragma omp critical
        if (DB[i] == key && i < position)
            position = i;
    }
}
```

a) Input BLOCK Geometric Data Decomposition:

- **Block size = DBsize / P**
- (DBsize % P) can be > 0
- Execution at each processor must end as soon as possible
- You cannot use *omp for*




```
#pragma omp parallel
{
    int myid = omp_get_thread_num ();
    int P = omp_get_num_threads ();
    unsigned long BS = DBsize / P;
    unsigned long lower = myid * BS;
    unsigned long upper = lower + BS;
    ...

    unsigned long i;
    for (i = lower; (i < upper ); i++) {
        #pragma omp critical
        if (DB[i] == key && i < position)
            position = i;
    }
}
```

a) Input BLOCK Geometric Data Decomposition:

- Block size = DBsize / P
- $(\text{DBsize} \% P)$ can be > 0
- **Execution at each processor must end as soon as possible**
- You cannot use *omp for*




```
#pragma omp parallel
{
    int myid = omp_get_thread_num ();
    int P = omp_get_num_threads ();
    unsigned long BS = DBsize / P;
    unsigned long lower = myid * BS;
    unsigned long upper = lower + BS;
    ...

    unsigned long i;
    for (i = lower; (i < upper && i < position); i++) {
        #pragma omp critical
        if (DB[i] == key && i < position)
            position = i;
    }
}
```

a) Input BLOCK Geometric Data Decomposition:

- Block size = DBsize / P
- **(DBsize % P) can be > 0**
- Execution at each processor must end as soon as possible
- You cannot use *omp for*

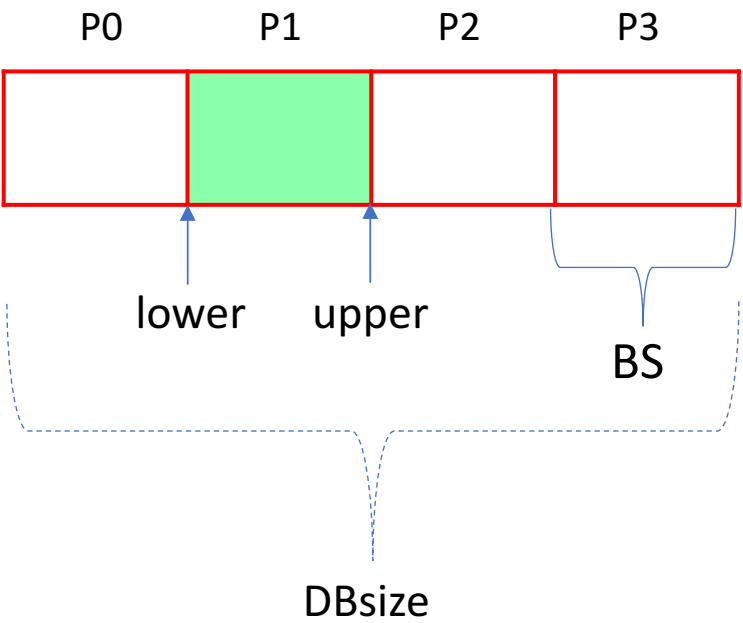


```
#pragma omp parallel
{
    int myid = omp_get_thread_num ();
    int P = omp_get_num_threads ();
    unsigned long BS = DBsize / P;
    unsigned long lower = myid * BS;
    unsigned long upper = lower + BS;
    unsigned long mod = DBsize % P;
    if (mod > 0)
        ...
}

unsigned long i;
for (i = lower; (i < upper && i < position); i++) {
    #pragma omp critical
    if (DB[i] == key && i < position)
        position = i;
}
}
```

a) Input BLOCK Geometric Data Decomposition:

- Block size = DBsize / P
- **(DBsize % P) can be > 0**
- Execution at each processor must end as soon as possible
- You cannot use *omp for*




```
#pragma omp parallel
{
    int myid = omp_get_thread_num ();
    int P = omp_get_num_threads ();
    unsigned long BS = DBsize / P;
    unsigned long lower = myid * BS;
    unsigned long upper = lower + BS;
    unsigned long mod = DBsize % P;
    if (mod > 0)
        ...
}
```

EJEMPLO:

		Total iterations = 15						
		Total number threads = 4						
	myid	iters						
	0	0	1	2				
	1	3	4	5				
	2	6	7	8				
	3	9	10	11	12	13	14	mod=3
	myid	iters						
id < mod	0	0	1	2	3			
	1	3	4	5	6	7		
	2	6	7	8	9	10	11	
id >= mod	3	9	10	11	12	13	14	x

a) Input BLOCK Geometric Data Decomposition:


- Block size = DBsize / P
- **(DBsize % P) can be > 0**
- Execution at each processor must end as soon as possible
- You cannot use *omp for*



```
#pragma omp parallel
{
    int myid = omp_get_thread_num ();
    int P = omp_get_num_threads ();
    unsigned long BS = DBsize / P;
    unsigned long lower = myid * BS;
    unsigned long upper = lower + BS;
    unsigned long mod = DBsize % P;
    if (mod > 0)
        if (myid < mod) {
            lower += myid; upper += myid + 1;
        } else {
            lower += mod; upper += mod;
        }
}


unsigned long i;
for (i = lower; (i < upper && i < position); i++) {
    #pragma omp critical
    if (DB[i] == key && i < position)
        position = i;
}
}
```


b) Optimize the code by reducing the serialization within the loop.



```
#pragma omp parallel
{
    int myid = omp_get_thread_num ();
    int P = omp_get_num_threads ();
    unsigned long BS = DBsize / P;
    unsigned long lower = myid * BS;
    unsigned long upper = lower + BS;
    unsigned long mod = DBsize % P;
    if (mod > 0)
        if (myid < mod) {
            lower += myid; upper += myid + 1;
        } else {
            lower += mod; upper += mod;
        }
}
unsigned long i;
for (i = lower; (i < upper && i < position); i++) {
    #pragma omp critical
        if (DB[i] == key && i < position)
            position = i;
}
```


b) Optimize the code by reducing the serialization within the loop.



```
#pragma omp parallel
{
    int myid = omp_get_thread_num ();
    int P = omp_get_num_threads ();
    unsigned long BS = DBsize / P;
    unsigned long lower = myid * BS;
    unsigned long upper = lower + BS;
    unsigned long mod = DBsize % P;
    if (mod > 0)
        if (myid < mod) {
            lower += myid; upper += myid + 1;
        } else {
            lower += mod; upper += mod;
        }
}

unsigned long i;
for (i = lower; (i < upper && i < position); i++) {
    #pragma omp critical
    if (DB[i] == key && i < position)
        position = i;
}
}
```

b) Optimize the code by reducing the serialization within the loop.




Filter the cases where it is known before entering the critical section that the if condition will evaluate to FALSE

```
#pragma omp parallel
{
    int myid = omp_get_thread_num ();
    int P = omp_get_num_threads ();
    unsigned long BS = DBsize / P;
    unsigned long lower = myid * BS;
    unsigned long upper = lower + BS;
    unsigned long mod = DBsize % P;
    if (mod > 0)
        if (myid < mod) {
            lower += myid; upper += myid + 1;
        } else {
            lower += mod; upper += mod;
        }
}

unsigned long i;
for (i = lower; (i < upper && i < position); i++) {
    #pragma omp critical
    if (DB[i] == key && i < position)
        position = i;
}
```

b) Optimize the code by reducing the serialization within the loop.



Filter the cases where it is known before entering the critical section that the if condition will evaluate to FALSE

```
#pragma omp parallel
{
    int myid = omp_get_thread_num ();
    int P = omp_get_num_threads ();
    unsigned long BS = DBsize / P;
    unsigned long lower = myid * BS;
    unsigned long upper = lower + BS;
    unsigned long mod = DBsize % P;
    if (mod > 0)
        if (myid < mod) {
            lower += myid; upper += myid + 1;
        } else {
            lower += mod; upper += mod;
        }
}

unsigned long i;
for (i = lower; (i < upper && i < position); i++) {
    if (DB[i] == key)
        #pragma omp critical
        if (i < position)
            position = i;
}
```


c) Write an alternative Geometric Data Decomposition that minimizes the average execution time.

```
#pragma omp parallel
{
    int myid = omp_get_thread_num ();
    int P = omp_get_num_threads ();
    unsigned long i;

    for (i = ... ; (i < ... && i < position); i ...) {
        if (DB[i] == key)
            #pragma omp critical
            if (i < position)
                position = i;
    }
}
```

Find as soon as possible the first position i that verifies $DB[i] == key$...

c) Write an alternative Geometric Data Decomposition that minimizes the average execution time.



```
#pragma omp parallel
{
    int myid = omp_get_thread_num ();
    int P = omp_get_num_threads ();
    unsigned long i;


    for (i = ...; (i < ... && i < position); i...) {
        if (DB[i] == key)
            #pragma omp critical
            if (i < position)
                position = i;
    }
}
```

Find as soon as possible the first position i that verifies $DB[i] == key$...



Data decomposition so that every thread has the same chance to find it:
Input Cyclic Geometric Data Decomposition

c) Write an alternative Geometric Data Decomposition that minimizes the average execution time.



```
#pragma omp parallel
{
    int myid = omp_get_thread_num ();
    int P = omp_get_num_threads ();
    unsigned long i;

    for (i = myid; (i < DBsize && i < position); i+= P) {
        if (DB[i] == key)
            #pragma omp critical
            if (i < position)
                position = i;
    }
}
```

Find as soon as possible the first position i that verifies $DB[i] == key$...



Data decomposition so that every thread has the same chance to find it:
Input Cyclic Geometric Data Decomposition